

← Go Back to Ensemble Techniques

☰ Course Content

Hands-On Quiz

Type	:	Practice Quiz
Attempts	:	1/1
Questions	:	10
Time	:	2h
Your Marks	:	18/20

Instructions

▼

Attempt History

Attempt #1

Jan 08, 9:55 PM

Marks: 18

▲

Q No: 1

Correct Answer

Marks: 2/2

Load the dataset and identify the variables that have more than 0.7 correlation with the dependent variable Opt_Service

A) WorkExp

B) Salary

C) MBA

D) Age

☐ Only A

☒ Both A and B

You Selected

☐ Both A and D

☐ A, B, C, and D

```
plt.figure(figsize=(10,7))
sns.heatmap(df.corr()[df.corr().>0.7],annot = True,vmax = 1, vmin = -1, cmap=
```

Q No: 2

Correct Answer

Marks: 2/2

Create dummies for the Gender column. Split the data into a 70:30 ratio. What is the percentage of 0 and 1 classes in the test data (y_test)?

Note - Do not use stratify on the dataset.



0: In a range of 0.8 to 0.95

You Selected

1: In a range of 0 to 0.1



0: In a range of 0.7 to 0.8

1 In a range of 0.2 to 0.3



0: In a range of 0.5 to 0.6

1: In a range of 0.3 to 0.6



0: In a range of 0.1 to 0.2

1: In a range of 0.3 to 0.5

```
# Creating X and Y variables
X = df.drop(['Opt_service'],axis=1)
y = df['Opt_service']
```

```
# creating dummies for Gender
X= pd.get_dummies(X, columns = ['Gender'],drop_first = True)
```

```
# Splitting data into training and test set:
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, ra
print(X_train.shape, X_test.shape)
```

```
#checking percentage of classes in y_test
print(y_test.value_counts(1))
```

Q No: 3

Incorrect Answer

Marks: 0/2

Build a bagging classifier with default parameters. How many employees who would take the service are correctly identified by the model from the training data?

☐ in a range of 10 - 20

☐ in a range of 20 - 30

Correct Option

☒ in a range of 30 - 40

You Selected

☐ in a range of 40 - 50

```
# fitting the baggin classifier on training set
bag = BaggingClassifier(random_state=1)
bag.fit(X_train, y_train)

# predicting on training set
y_train_pred = bag.predict(X_train)

# creating the confusion matrix
sns.heatmap(confusion_matrix(y_train,y_train_pred),annot=True,fmt='.1f')
plt.xlabel('Predicted values')
plt.ylabel('Actual values')
plt.show()
```

Q No: 4

Correct Answer

Marks: 2/2

Build a random forest classifier with default parameters and a bagging classifier with logistic regression as the base estimator.

Select which of the following statements is/are true?

- A) Random forest is giving the same recall as bagging classifier on the train data
- B) Random forest is giving a higher recall than bagging classifier on the train data
- C) Random forest is giving a higher f1_score than bagging classifier on the train data
- D) Random forest is giving a lower f1_score than bagging classifier on the train data

☐ A and D

☐ C and D

☐ A and C

☒ B and C

You Selected

```
# Random Forest
rf = RandomForestClassifier(random_state=1)
rf.fit(X_train, y_train)

# bagging classifier with base estimator =Logistic Regression
bag_lr = BaggingClassifier(random_state=1, base_estimator = LogisticRegression)
bag_lr.fit(X_train, y_train)

# predicting on the training set
rf_predict = rf.predict(X_train)
bag_predict = bag_lr.predict(X_train)

# Performance of Random Forest

print('Random Forest')

print(recall_score(y_train, rf_predict))
print(f1_score(y_train, rf_predict))

# Performance of Bagging classifier

print('Bagging')

print(recall_score(y_train, bag_predict))
print(f1_score(y_train, bag_predict))
```

Q No: 5

Correct Answer

Marks: 2/2

Build a bagging classifier model with the base estimator as a decision tree.

Vary the depth of the base estimator/Decision tree from depth 1 to 5 (both values included) and compare their performance.

Select which of the following statements is true?

- A) At depth = 1, the bagging classifier gives the highest f1 score among all the models on the training set.
- B) At depth = 2, the bagging classifier gives the highest f1 score among all the models on the training set.
- C) At depth = 4, the bagging classifier gives the highest f1 score among all the models on the training set.
- D) At depth = 5, the bagging classifier gives the highest f1 score among all the models on the training set.

☐ Only A

☐ Only B

☐ Only C

☒ Only D

You Selected

```
scores = []
for i in range(1,6):
    bag = BaggingClassifier(base_estimator=DecisionTreeClassifier(random_state=
    bag.fit(X_train, y_train)
    pred = bag.predict(X_train)
    case = {'Depth':i, 'F1 Score':f1_score(y_train,pred)}
    scores.append(case)
```

```
print(scores)
```

Q No: 6

Correct Answer

Marks: 2/2

Which of the boosting models (AdaBoost, GradientBoost, XGBoost) with default parameters has the lowest f1-score on the training set?

Note: Set the eval_metric = 'logloss' for XGBoostClassifier

☐ AdaBoost

☐ Gradient boosting

☐ Xgboost

☒ All three models have the same f1 score

You Selected

```
# Adaboost classifier
abc = AdaBoostClassifier(random_state=1)
abc.fit(X_train, y_train)
```

```
# GradientBoost classifier
gbm = GradientBoostingClassifier(random_state=1)
gbm.fit(X_train, y_train)
```

```
#XGBoost classifier
xgb = XGBClassifier(random_state=1, eval_metric = 'logloss')
xgb.fit(X_train, y_train)
```

```
# predicting on training set
abc_predict = abc.predict(X_train)
gbm_predict = gbm.predict(X_train)
xgb_predict = xgb.predict(X_train)
```

```
# Checking model performance
print('ABC')
print(f1_score(y_train, abc_predict))
```

```
print('GBM')
print(f1_score(y_train, gbm_predict))
```

```
print('XGB')
print(f1_score(y_train, xgb_predict))
```

Q No: 7

Correct Answer

Marks: 2/2

Use the gradient boosting classifier trained in Q6 and plot the feature importance of the variables. Which are the 2 most important variables respectively?

☐ MBA and Engineer

☒ Salary and Distance

You Selected

☐ Age and WorkExp

☐ License and Gender_Male

```
# Gradient boosting model
gbm = GradientBoostingClassifier(random_state = 1)
gbm.fit(X_train, y_train)
```

```
# visualizing feature importance
feature_names = X_train.columns
importances = gbm.feature_importances_
indices = np.argsort(importances)
plt.figure(figsize=(12, 12))
plt.title("Feature Importances")
plt.barh(range(len(indices)), importances[indices], color="violet", align="center")
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel("Relative Importance")
plt.show()
```

Q No: 8

Correct Answer

Marks: 2/2

Train three models:

1. Model1 = Gradient Boosting classifier with $n_{\text{estimator}} = 50$ and learning rate = 0.01
2. Model2 = Gradient Boosting classifier with $n_{\text{estimator}} = 100$ and learning rate = 0.01
3. Model3 = Gradient Boosting classifier with $n_{\text{estimator}} = 400$ and learning rate = 0.01

f1_score1, f1_score2, f1_score3 are f1_scores of the three models respectively.

What is the order of f1_score on the training set?

☐ f1_score1 > f1_score2 > f1_score3

☐ f1_score2 > f1_score1 > f1_score3

☒ f1_score3 > f1_score2 > f1_score1

You Selected

☐ f1_score3 > f1_score1 > f1_score2

```
# Training first model
gbm1 = GradientBoostingClassifier(random_state=1, n_estimators=50, learning_rate=0.1)
gbm1.fit(X_train,y_train)

# Training second model
gbm2 = GradientBoostingClassifier(random_state=1, n_estimators=100, learning_rate=0.1)
gbm2.fit(X_train,y_train)

# Training third model
gbm3 = GradientBoostingClassifier(random_state=1, n_estimators=400, learning_rate=0.1)
gbm3.fit(X_train,y_train)

# Checking model performance
print('First Model Result')
y_pred = gbm1.predict(X_train)
print(f1_score(y_train,y_pred))

print('Second Model Result')
y_pred = gbm2.predict(X_train)
print(f1_score(y_train,y_pred))

print('Third Model Result')
y_pred = gbm3.predict(X_train)
print(f1_score(y_train,y_pred))
```

Q No: 9

Correct Answer

Marks: 2/2

Build a stacking classifier using two models - Decision Tree, Bagging Classifier as base estimators and use Random Forest as the final estimator. Which of the following is/are true?

Select which of the following statements is/are true?

- A) Stacking classifier is giving f1 score in the range of 0.7 - 0.9 on the training set
- B) Stacking classifier is giving recall score in the range of 0.8 - 1 on the training set
- C) Stacking classifier is giving f1 score in the range of 0.5 - 0.7 on the training set
- D) Stacking classifier is giving recall score in the range of 0.65 - 0.75 on the training set

☒ A and B

You Selected

☐ B and C

☐ C and D

☐ A and D

```
# creating list of estimators
estimators = [('DT',DecisionTreeClassifier(random_state=1)),
              ('BAG',BaggingClassifier(random_state=1))]

# Training stacking classifier
clf = StackingClassifier(estimators = estimators, final_estimator = RandomForestClassifier())
clf.fit(X_train, y_train)

# Predicting on training set
y_pred = clf.predict(X_train)

# Checking model performance
print('F1 score')
f1_score(y_train,y_pred)

print('Recall score')
recall_score(y_train,y_pred)
```

< Previous

Next >

